

Technical Manual of the Ceylan-Myriad Layer



Organisation: Copyright (C) 2008-2021 Olivier Boudeville

Contact: about (dash) myriad (at) esperide (dot) com

Creation date: Sunday, August 17, 2008

Lastly updated: Friday, August 6, 2021

Version: 1.0.19

Status: Stable

Dedication: Users and maintainers of the Myriad layer.

Abstract: The role of the [Myriad](#) layer (part of the [Ceylan](#) project) is to gather all [Erlang](#) general-purpose base constructs that we found useful for (Erlang-based) developments.

We present here a short overview of these services, to introduce them to newcomers. The next level of information is either to browse the [Myriad API documentation](#) or simply to read the corresponding [source files](#), which are intensely commented and generally straightforward.

The latest version of this documentation is to be found at the [official Ceylan-Myriad website](#) (<http://myriad.esperide.org>).

This documentation is also mirrored [here](#).

Table of Contents

Overview & Context	4
Usage Guidelines	4
License	4
About Layers	5
Recommended Usage & Contribution	5
Getting Myriad	5
Prerequisites	5
Getting Myriad's Sources	6
Building Myriad	6
Testing Myriad	7
Type-checking Myriad	7
Maintaining Myriad and Deriving Projects with regard to rebar3	7
OTP Build	8
Why Providing Two Different Build/Deploy/Run Systems	8
Relying on Rebar3	9
OTP Application	9
OTP Release	11
Hex Package	11
Other OTP-related Make Targets of Interest	12
Services offered by the Myriad Layer	13
General Build Structure	14
General Settings	15
Maths Services	16
Data-Management Services	18
Datatypes	18
File Formats	19
Support for Code Injection	21
Defining a token	21
Defining the code to inject	21
Using tokens to enable code injection	22
Controlling assertions	24
Usage Hints	25
For more information	25
Helpers For User Interface Programming	26
Various Flavours of User Interfaces	26
Raw Text User Interface: <code>text_ui</code>	27
Terminal Text User Interface: <code>term_ui</code>	27
Graphical User Interface: <code>gui</code>	27
Myriad's Helper Shell Scripts	29
Erlang-Dedicated Scripts	29
More General Scripts	32
Script-based Apps	33
Utility Toolbox	36
Support for Metaprogramming	39
Management of Units	41

Motivation	41
Available Support	41
Specifying Units	41
Possible Improvements Regarding Dimensional Analysis	44
SQL support	46
About SQL	46
Database Back-end	46
Erlang SQL Binding	46
SQL Support Provided By the <i>Myriad</i> Layer	48
SQL-related Troubleshooting	48
Myriad Main Conventions	49
Text Conventions	49
Coding Practices	49
Execution Targets	53
Geometric Conventions	53
Other Conventions	54
Myriad-related Troubleshooting	55
Header/Module Dependencies	55
Third-Party Dependencies	55
Runtime-only Third-Party Dependencies	55
Build-time Third-Party Dependencies	56
Myriad-level Third-Party Dependencies	56
About the <code>table</code> module	57
Enabling the Interconnection of Erlang nodes	57
Safety Measures	57
Testing & Troubleshooting	58
Support for Myriad	59
Please React!	60
Contributions & Ending Word	60

Overview & Context

When using any programming language, there are always **recurring patterns** that prove useful.

Instead of writing them again and again, we preferred gathering them all in a **low-level layer** (mostly a modest **code library**), in their most convenient, reliable, efficient version, together with their specification, documentation and testing.

This layer provides its (generally lightweight, simple) services just on top of the [Erlang](#)³ language, as a relatively small (comprising currently about 70k lines), thin layer.

These services tend to stay away from introducing any new dependency.

Should a key, generic service need a third-party prerequisite (ex: library to manage a complex data format, or to process specific data), that dependency should be made fully optional (see the [Third-Party Dependencies](#) section for more details).

As a consequence, for the [Ceylan](#) project, the first level of the Erlang-based software stack that we use relies on this Myriad layer - whose official, more specific name is the [Ceylan-Myriad](#)⁴ layer.

The project repository is located [here](#).

Usage Guidelines

License

The Myriad layer is licensed by its author (Olivier Boudeville) under a disjunctive tri-license, giving you the choice of one of the three following sets of free software/open source licensing terms:

- the [Mozilla Public License](#) (MPL), version 1.1 or later (very close to the former [Erlang Public License](#), except aspects regarding Ericsson and/or the Swedish law)
- the [GNU General Public License](#) (GPL), version 3.0 or later
- the [GNU Lesser General Public License](#) (LGPL), version 3.0 or later

This allows the use of the Myriad code in as wide a variety of software projects as possible, while still maintaining copyleft on this code.

Being triple-licensed means that someone (the licensee) who modifies and/or distributes it can choose which of the available sets of licence terms he/she is operating under.

Enhancements are expected to be back-contributed (hopefully), so that everyone can benefit from them.

³If needing to discover/learn Erlang, we recommend browsing [Learn You Some Erlang for great good!](#) or, even better, buying their book!

⁴It was formerly known as the **Common** layer.

About Layers

The **Myriad** services are to be used by this layer itself (for its inner workings), and, more importantly, are to be re-used, specialised and enriched by layers built on top of it.

The general rule is that a layer may depend on (i.e. make use of) all layers *below* it (not only the one just preceding it), but cannot refer to any layer *above* it (it should be literally unaware of their existence).

So, in a bottom-up view, generally a software stack mentioned here begins with the operating system (typically GNU/Linux), then [Erlang/OTP](#), then **Myriad**, then any layer(s) built on top of them (ex: [WOOPER](#)).

Of course a given layer does not mask the layers below; for example programs using the **Myriad** layer typically use also a lot the services brought by the [Erlang base libraries](#).

Recommended Usage & Contribution

When developing Ceylan-based code, if needing a service already provided by this **Myriad** layer, it is strongly advised to use that service and, possibly (if useful), expand or enrich it, with backward compatibility in mind.

If such a service is not provided by the current version of the layer, yet is deemed generic enough, then it should preferably be added directly to the relevant part of the library and called from the code that was needing it.

Of course, contributions of all sorts are welcome.

We do our best to test, at least lightly, each element provided. All services offered in a `foo.erl` file are thus expected to be tested in the companion `foo_test.erl` file, in the `test` tree (whose structure tends to mirror the one of the `src` tree). Once there, running this test is as simple as executing:

```
$ make foo_run
```

Note that however we have not reached the discipline level of an exhaustive `eunit` test suite for each service (most of them being almost trivial).

The [Dialyzer](#) static analysis tool is regularly run on the code base (see the `generate-local-plt` and `self-check-against-plt` generic Make targets for that).

Getting Myriad

Prerequisites

The **operating system** is supposed to be any not-so-old GNU/Linux distribution⁵.

People reported uses of **Myriad** on **macOS**, yet no extensive testing has been done there.

Whereas Erlang supports **Windows** and we tried to be as cross-platform as possible, even with tools like **MSYS2** / **MinGW-w64** we suppose that quite a lot of special cases would have to be addressed (patches welcome, though!).

⁵For what it is worth, we prefer [Arch Linux](#), but this does not really matter here.

The main tool prerequisite is of course having the [Erlang](#) environment available, in its 24.0 version⁶ or more recent. There are various ways of obtaining it (from your distribution, from prebuilt packages, directly from the sources), one of which being the [install-erlang.sh](#) script that we devised. A simple use of it is:

```
$ ./install-erlang.sh --doc-install --generate-plt
```

One may execute `./install-erlang.sh --help` for more guidance about how to configure it, notably in order to enable all modules of interest (`crypto`, `wx`, etc.).

Getting Myriad's Sources

This is pretty straightforward, based on the [project repository](#) hosted by Github:

```
$ git clone https://github.com/Olivier-Boudeville/Ceylan-Myriad.git myriad
```

This should download in your current directory the full Myriad repository. For OTP compliance, using for such a clone its short name (`myriad`) rather than its long one (`Ceylan-Myriad`) is recommended.

The Myriad `master` branch is meant to stick to the latest stable version: we try to ensure that this main line always stays functional (sorry for the pun). Evolutions are to take place in feature branches and to be merged only when ready.

Building Myriad

If a relevant Erlang installation is available, this is as simple as:

```
$ cd myriad
$ make all
```

The parallel build of the whole layer (services and tests alike) shall complete successfully (if it is not the case, see our [support](#) section).

One may just run `make` by itself in order to list the main available options.

Note that by default our native, make-based, build system is used. Alternatively, a `rebar3`-based build can be done (see the [OTP Build](#) section for more details).

In this case one may run `make create-myriad-checkout` in order to create, based on our conventions, a suitable `_checkouts` directory so that `rebar3` can directly take into account local, directly available (in-development) dependencies (although Myriad does not have any, beside Erlang itself - this `make` target is useful for the layers built on top of Myriad).

So, alternatively to using `make` directly, one may execute `rebar3 compile` instead.

⁶Most probably that older versions of Erlang would be more than sufficient in order to build Myriad (possibly at the expense of minor changes in a few calls to standard modules having been deprecated since then). It is just that in general we prefer to stick to the latest stable versions of software such as Erlang, and advise you to do so.

To determine the recommended version of Myriad-based code, just execute our [install-erlang.sh](#) script with its `--version` option.

Testing Myriad

As Myriad has no prerequisite (besides Erlang itself of course), just run (possibly simply thanks to `rebar3 compile` after a `git clone https://github.com/Olivier-Boudeville/Ceylan-Myriad.git`), still from the root directory of Myriad:

```
$ make test
```

The testing shall complete successfully (if it is not the case, see our [support](#) section).

Note

Myriad is built and tested at each commit through [continuous integration](#). The same holds for the projects based on it, directly (ex: [WOOPER](#), [Seaplus](#)) or not (ex: [Traces](#), [Mobile](#), [US-Web](#)), so in terms of usability, confidence should be high.

Type-checking Myriad

As Myriad is (by default) to enable debug information with a key-based protection of the resulting BEAM files, one should first have such key defined.

One way of doing so is, if wanted, to update the default key (see `DEBUG_INFO_KEY` in `GNUmakevars.inc`) and to write in on disk (ex: `make write-debug-key-file`), and to rebuild Myriad accordingly afterwards (ex: `make rebuild`).

Then, still from the `myriad` root directory:

```
$ make generate-local-plt self-check-against-plt
```

It will trigger a full type-checking of Myriad, done thanks to [Dialyzer](#).

This time-consuming phase will complete with a rather long list of notifications. Help us reducing it! These messages are numerous, but we do not think that most of them are so worrying.

Finally, to trigger in one go a full rebuild, testing and type-checking, one may run:

```
$ make check
```

Maintaining Myriad and Deriving Projects with regard to `rebar3`

For Myriad as for all developments built on top of it (ex: specialisation layers or applications), any dependency may be specified in their `rebar.config`⁷ through a branch of a GIT repository corresponding to that dependency.

For example, Myriad itself does not require any specific dependency, but projects making use of Myriad (ex: [WOOPER](#)) may specify in their `rebar.config`:

⁷For example, with the conventions we rely on, `rebar.config` is generated from the `conf/rebar.config.template` file of the project of interest.

```
{deps, [{myriad, {git, "git://github.com/Olivier-Boudeville/Ceylan-Myriad",
                    {branch, "master"}}}]}
```

However, when having to build a dependency, rebar3 will not refer to the tip of the branch specified for it, but to any commit it may read from any pre-existing `rebar.lock` file at the root of the current project (the underlying goal being to allow for more reproducible builds).

As the [rebar3 recommendation](#) is to store a version of that lock file in source version control, **it shall be regularly updated** otherwise the dependencies of a given project will stick, for the worst, to an older version of their branch, designated by an obsolete reference (this can be detected for example when continuous integration breaks after a nevertheless legit commit of the project). The solution is thus, for a project of interest, to regularly force an update of its dependencies referenced in its own lock file, and to commit the resulting version.

For example, one would issue from the root of the project of interest:

```
$ rebar3 upgrade
```

This may update the `ref` entry of its dependencies (including Myriad) in its `rebar.lock` file, which shall then be committed for posterity.

OTP Build

These build considerations apply to Myriad but also, more generally, to most if not all our Erlang developments.

Why Providing Two Different Build/Deploy/Run Systems

We felt that OTP build tools and Emakefiles were not expressive enough for our needs: as mentioned in [Building Myriad](#), a full, rather complete/complex/powerful build system based on [GNU make](#) is used by Ceylan-Myriad natively instead, and has been fully satisfactory for years (simple, lightweight, reliable, controllable, flexible, fast, etc.).

It allows to introduce all the generic rules we wanted, to define many conditional settings, to walk through an arbitrarily nested source tree, to integrate within a layered stack (notably alongside some other **Ceylan-*** libraries that depend on Ceylan-Myriad) and to perform a multi-stage build to accommodate the compilation and use of parse-transforms, with their own set of prerequisites.

More precisely we routinely (see [WOOPER](#) or [Seaplus](#)) rely on layers built on top of Myriad, which define their own parse transforms that are themselves parse-transformed by Myriad's one - and it works great.

However, to better integrate with other Erlang developments (which are mostly OTP-compliant), we added the (optional) possibility of generating a Myriad *OTP library application* out of the build tree, ready to be integrated into an (OTP) *release* and to be available as an Hex *package*. For that we rely on [rebar3](#), [relx](#) and [hex](#).

So currently all our Erlang-based developments can also be built and tested through rebar3, and this support is checked at each commit thanks to continuous integration.

We use less frequently releases (we rely on a basic deployment procedure of our own) and even less hex, yet they were supported once, so we believe that their integration should be at least fairly close to be operational (if not, patches welcome!).

Relying on Rebar3

Despite the kind support of the rebar3 authors and much time spent on its integration, sometimes our build based on it (for Myriad and the projects built on top of it) has encountered issues or has been lagging behind our native one.

Now we believe that all pending issues have been solved (rebar3 is a neat tool), yet being able to switch back to another lighter, ad-hoc, more controlled build system is sometimes a relief - at least a welcome security. Anyway the user can choose between these two (native vs rebar3) build machineries. As for us, we still prefer our native build system, even if it leaves to the developer the task of installing the needed prerequisites by him/herself.

OTP Application

Myriad is not an *active* OTP application, and as such does not rely on, or provides, services running in the background; so no supervision tree or `gen_server` is involved here, just a *library* application ready for OTP integration⁸.

Getting rebar3 There are [various ways](#) for obtaining `rebar3`; we prefer:

```
$ cd ~/Software && git clone https://github.com/erlang/rebar3.git
&& cd rebar3 && ./bootstrap
```

Alternatively, should you just want to update a (pre-existing) rebar3 install, first get the current version (`rebar3 -v`) to check it afterwards, then issue `rebar3 local upgrade`; however this would involve running rebar from `.cache/rebar3/bin`, so instead we prefer using (typically from `~/Software/rebar3`):

```
$ git pull && ./bootstrap
```

Another option is to download a prebuilt version of rebar3.

Finally, one may prefer using the [install-rebar3.sh](#) script that we devised, which automates and enforces our conventions while letting the choice between an installation from sources or from a prebuilt version thereof (just un `install-rebar3.sh --help` for guidance).

⁸Speaking of OTP, in development mode, `proc_lib`-based spawns used to be enabled, yet this led to longer error messages that were not that useful; see `spawn_utils.hrl` if wanting to re-enable them.

Generating Ceylan-Myriad Then, from the root of a Myriad clone, to obtain the Ceylan-Myriad library *application*, one just has to enter:

```
$ make rebar3-application
```

It will trigger `rebar3`, resulting in a full, OTP-compliant build tree created in `_build` (including a properly-generated `_build/default/lib/myriad/ebin/myriad.app` file), and more generally in a proper OTP application⁹.

A full, autonomous, functional by design build procedure can be also found in Myriad's [continuous integration](#) script.

Testing Ceylan-Myriad As a result, the OTP application support can be tested from the root of an (already-built, with `make rebar3-application`) Myriad source tree:

```
$ cd src/utils
$ make myriad_otp_application_run
Running unitary test myriad_otp_application_run (third form) from
myriad_otp_application_test

--> Testing module myriad_otp_application_test.

Starting the Myriad application.
Myriad version: {1,0,11}.
Current user name: 'stallone'.
Stopping the Myriad application.
Successful end of test of the Myriad application.
=INFO REPORT==== 18-Jul-2019::22:37:24.779037 ===
    application: myriad
    exited: stopped
    type: temporary

--> Successful end of test.

(test finished, interpreter halted)
```

This support can be also tested manually, directly through the build tree used by `rebar3`; from the root of Myriad, after having run `make rebar3-application`:

```
$ erl -pz _build/default/lib/myriad/ebin/
Erlang/OTP 22 [erts-10.4] [source] [64-bit] [smp:8:8] [...]

Eshell V10.4 (abort with ^G)
1> application:start(myriad).
ok
```

⁹The rebar-based build relies, thanks to `{pre,post}-compile` hooks, on our native build system. Because of extraneous, failing recompilations being nevertheless triggered by rebar, we had to introduce bullet-proof hooks (refer to [1](#), [2](#)).

```

2> text_utils:format( "Hello ~s", [ world ] ).
"Hello world"
3> application:stop(myriad).
=INFO REPORT==== 18-Jul-2019::22:47:36.429804 ===
    application: myriad
    exited: stopped
    type: temporary

```

When needing to include a Myriad header file (taking `spawn_utils.hrl` as an example) in one's code, OTP conventions mandate using:

```

-include_lib("myriad/include/spawn_utils.hrl").

rather than:

-include("spawn_utils.hrl").

```

OTP Release

Quite similarly, to obtain a Ceylan-Myriad OTP *release* ([relx](#) being used in the background), possibly for a given profile like `default` (development mode) or `prod` (production mode) - refer to `REBAR_PROFILE` in `GNUmakevars.inc`, one just has to run, from the root of Myriad:

```
$ make rebar3-release
```

Hex Package

The [hex](#) package manager relies on `mix`, which is commonly installed with [Elixir](#) (another language built on top of the Erlang VM).

Thanks to the `rebar3` integration with the `rebar3_hex` plugin specified in Myriad's (generated) [rebar.config](#), `hex` will be automatically installed and set up.

By following the publishing guidelines ([\[1\]](#), [\[2\]](#)), we were able to publish [Hex packages for Myriad](#) that can be freely used. And there was much rejoicing!

One just has to specify for example `{deps, [myriad]}` in one's `rebar.config`, and that's it.

Note

Finally our workflow does not rely on Hex, so we do not update the Hex packages anymore. Just drop us an email if needing an updated one.

For more details, one may have a look at:

- [rebar.config.template](#), the general rebar configuration file used when generating the Myriad OTP application and release
- [rebar-for-hex.config.template](#), to generate a corresponding Hex package for Myriad (whose structure and conventions is quite different from the previous OTP elements)
- [rebar-for-testing.config.template](#), the simplest test of the previous Hex package: an empty rebar project having for sole dependency that Hex package

Other OTP-related Make Targets of Interest

To populate/update the OTP build tree (by default, from the GIT root, for example `_build/default/lib/myriad/` for Myriad) of the current Ceylan layer, one may use:

```
$ make rebar3-compile
```

(this is especially useful in order to be able to use directly, from an OTP application, changes just performed in a Ceylan-based layer)

To update both the OTP build tree and the local ebin directory of each Ceylan layer on which the current layer depends, use:

```
$ make rebar3-local-update
```

(note this will be a no-op from Myriad, as it does not depend on any Ceylan layer)

To publish an Hex package (once the proper version number has been set in `GNUmakevars.inc`, see `MYRIAD_VERSION`):

```
$ make rebar3-hex-publish
```

To test such a package:

```
$ make test-hex-package
```

To populate directly the OTP local build tree with the Ceylan dependencies located alongside the current install (not useful for Myriad - which depends on none, but useful for upper layers) rather than fetching them through Hex (otherwise many more Hex packages would have to be published for testing during development):

```
$ make rebar3-local-update
```

Many more targets are defined in [GNUmakrules-explicit.inc](#).

Services offered by the Myriad Layer

The **Myriad** services are gathered into following themes:

1. General [build structure](#)
2. General [settings](#)
3. [Maths](#) services
4. [Data-management](#) services
5. Helpers for [graphical user interface](#) (GUI) programming
6. All-purpose [helper scripts](#)
7. Utility [toolbox](#)
8. Management of [units](#)
9. [Metaprogramming](#), based on heavy use of parse transforms
10. [SQL support](#)

In future versions of this document, following topics will be discussed:

- HDF5 support
- REST support
- third-party language bindings (ex: Python, Java, maybe in the future Haskell; C/C++ is to be tackled by our [Seaplus](#) project)
- RDF support

Even if this document does not constitute an exhaustive walk-through, each of them is detailed in turn below.

The next level of detail is either to browse the [Myriad API documentation](#) or simply to peer at the referenced [source files](#), as they include many implementation notes, comments and typing information.

General Build Structure

Various elements are defined at the `myriad` level to set-up an appropriate build, based on [GNU Make](#).

This includes:

- a set of pre-defined Make **variables**, describing various settings that will be reused by generic rules (ex: to compile modules with relevant flags, to create source archives, to install an application, to manage the various paths, to perform test checking, to generate archives, installs and releases, etc.); these variables are defined in `myriad/GNUmakevars.inc`
- a set of generic **rules**, to compile and run various modules and tests, to generate various elements of documentation, etc.; these rules are defined (still from the `myriad` root directory), in:
 - `GNUmakerules-automatic.inc`, for all rules that apply generically to a some kinds of targets
 - `GNUmakerules-explicit.inc`, for all "direct" rules, that are not pattern-based
 - `doc/GNUmakerules-docutils.inc`, for all documentation-related rules
- finally, the whole is gathered in a unique file to include, `GNUmakesettings.inc`, whose structure allows for a safe and sound combination of all these build element across a series of layers (the first of which being `Myriad`)
- **examples** of minimal Make files, which mostly specify the relative base path and only refer to the generic variables and rules; see `myriad/src/GNUmakefile` as an example

An example of this stacked structure is the `Ceylan-WOOPER` layer (see [official site](#)), which is directly built on top of `Ceylan-Myriad` (and itself a base layer for other layers and applications).

These build facilities are designed to be enriched in turn by all layers above, which may add or override variables and rules.

As an example, the `shell` make target allows to spawn an Erlang shell that is readily able to use the current layer and all the ones below (including thus `Myriad`); this way direct, hassle-free interactive testing can be performed - which may prove quite convenient.

General Settings

These general-purpose settings and helpers, gathered in Myriad's [conf](#) directory, deal with:

- default CSS files (`Default-docutils.css`)
- our recommended versions of (commented) configuration files for various tools:
 - for Emacs: `init.el`, to be placed in the `~/.emacs.d/` directory; in addition to our `init.el` configuration file, we rely on `acme-search.el`, `flyspell-guess.el` and `whitespace.el` (all in the `~/.emacs.d` directory - check that they are not eclipsed by any other initialisation file, like a stray `~/.emacs` file; see also our [update-emacs-modules.sh](#) script)
 - for Nedit: `nedit.rc`, to be placed in the `~/.nedit/` directory
- our standard script to properly install Erlang ([install-erlang.sh](#)) with detailed comments and command-line options (use `install-erlang.sh --help` for more information)

Maths Services

Some simple maths-related operations are defined in the `myriad/src/maths` directory:

- the most basic services are centralised in `math_utils.erl` and provide:
 - **general operations** apparently lacking to Erlang (for example for conversions or rounding (`floor/1`, `ceiling/1`), or not exactly implemented as we would have liked (ex: `modulo/2`)
 - operations tailored to operate on **floating-point values** (ex: `are_close/2,3`, `are_relatively_close/2,3`, `get_relative_difference/2`, `is_null/1`)
 - operations on **angles** (ex: `radian_to_degree/1`, `canonify/1`)
 - the associated **typing** information
- linear-related operations are defined; for example the **2D** operations are defined in `linear_2D.erl` (their **3D** counterparts being defined in `linear_3D.erl`, their **4D** counterparts in `linear_4D.erl`; base ones in `linear.erl`) and include:
 - operations on **points**: `are_close/2`, `is_within/3`, `square_distance/2`, `distance/2`, `cross_product/2`, `roundify/1`, `get_integer_center/2`, `get_center/2`, `translate/2`, etc.
 - operations on **vectors**: `vectorize/2`, `square_magnitude/1`, `magnitude/1`, `scale/2`, `make_unit/1`, `normal_left/1`, `normal_right/1`, `dot_product/2`, etc.
 - operations on **lines**: `get_line/2`, `intersect/2`, `get_abscissa_for_ordinate/2`, etc.
 - operations related to **angles**: `is_strictly_on_the_right/3`, `is_obtuse/1`, `abs_angle_rad/3`, `angle_rad/3`, `abs_angle_deg/3`, `angle_deg/3`, etc.
 - operations on **sets of points**: `compute_smallest_enclosing_rectangle/1`, `compute_max_overall_distance/1`, `compute_convex_hull/1`, etc.
- **polygon**-related operations are available in `polygon.erl`:
 - **generation** of polygons: `get_triangle/3`, `get_upright_square/2`, `get_polygon/1`, etc.
 - **operations** on them: `get_diameter/1`, `get_smallest_enclosing_rectangle/1`, `get_area/1`, `is_in_clockwise_order/1`, `is_convex/1`, `to_string/1`, etc.
 - **rendering** them: `render/2`, `set_edge_color/2`, `get_edge_color/1`, `set_fill_color/2`, `get_fill_color/1`, etc.
 - managing their **bounding boxes**: `update_bounding_box/2`, etc.
- **bounding-boxes in general** are supported in `bounding_box.erl`, including `get_lazy_circle_box/1`, `get_minimal_enclosing_circle_box/1`, etc.

- a minimalist [Runge-Kutta solver](#) is defined in `rk4_solver.erl`

A few geometry-related conventions used in Myriad:

- implicit homogeneous w coordinates are 1.0
- front-facing polygons are the ones whose signed area is strictly positive: their vertices shall be seen in counter-clockwise winding order when front-facing towards the viewer (as in OpenGL's `GL_CCW`); see also: `polygon:{get_area,get_signed_area}/1`

Data-Management Services

Datatypes

Some generic **data-structures**, in addition to the ones provided built-in with Erlang, are defined in `myriad/src/data-management`, and described below.

Table Types A set of **associative tables** is available, each offering a rather complete interface (to create, update, enrich, search, query, list, map, fold, merge, display, etc. a table or entries thereof) and a different trade-off. Various implementations are defined (with tests and benchmarks), in:

```
{hash, lazy_hash, list_, tracked_hash, map_hash}table.erl
```

A **table pseudo-module** is additionally provided, in order to abstract out these various options: the user may then rely exclusively on `table`, regardless of the actual table type to which it will be translated at compilation-time.

Indeed the `table` module is a fully virtual one, in the sense that neither `table.erl` nor `table.beam` exists, and that the Myriad parse transform is to automatically replace a call to this `table` pseudo-module by a call to one of the aforementioned table types (and it will do the same replacement for type specifications as well).

By default, such `table` calls will be translated to corresponding calls to the `map_hashtable` module, which is generally the most efficient one (it relies on the more recently-introduced Erlang built-in `maps` datatype, which `table` now favors).

As a result, in order to consult the `table` API, please refer to `map_hashtable.erl`.

Such a default implementation may be overridden on a per-module basis, thanks to a `table_type` define.

For example, specifying `-table_type(list_table).` will result in the current module to translate `table` to `list_table`, instead of the default `map_hashtable`.

Finally, a way of **generating a read-only associative table** whose key/value pairs can be read very efficiently from any number (potentially extremely large) of readers (processes) is provided with `const_table.erl`.

No ETS table, replication (ex: per-user table copy) or message sending is involved: thanks to meta-programming, a module is generated on-the-fly, exporting as many functions as there are different keys in the table of interest.

Calling a function corresponding to a key returns the associated value.

More precisely, a module name (ex: `foobar`) and a `table:table(atom(), any())` instance shall be provided to `const_table:generate/2`; then, for each key/value pair in the specified table (ex: `{baz, 42.0}`), a 0-arity function is generated and exported in that module, as if we had:

```
-module(foobar).  
  
[...]  
  
-export([baz/0]).
```

```
-spec baz() -> term().
baz() ->
    42.0.
```

Then third-party code can call for example `foobar:foo()` and have `42.0` returned. This is presumably the most efficient way of sharing constants in Erlang.

Note that no actual module file is generated (ex: no `foobar.beam` written in the filesystem), the operation remains fully in-memory.

Keys must be atoms, and the table of interest shall be immutable (`const`), even if, thanks to hot code upgrade, one may imagine updating the table at will, having any number of successive versions of it.

Generating a table of the same name more than once should be done with care, as if a given table is generated thrice, the initial table would first switch from "current" to "old", and then would be removed. Any process that would linger in a function of this module would then be terminated (see [code replacement](#)).

However, due to the nature of these tables (just one-shot fully-qualified calls, no recursion or message-waiting construct), this is not expected to happen.

Other Datatypes They include `pair.erl`, `option_list.erl`, `preferences.erl`, `tree.erl`.

One may also refer to `set_utils.erl` and `ring_utils.erl`.

Pseudo-Builtin Types Such types, as `void/0`, `maybe/1` and `fallible/{1,2}` are supported, thanks to the Myriad parse-transform.

File Formats

Basic File Formats A built-in very basic support for the [CSV](#) (see `csv_utils`) and [RDF](#) (see `rdf_utils`) conventions is provided.

Most Usual, Standard File Formats An optional support (as it depends on third-party prerequisites) is proposed for:

- JSON
- HDF5
- SQLite

Refer to the [Myriad-level Third-Party Dependencies](#) section for further information.

For Pure Erlang uses: ETF File Format For many needs in terms of Erlang internal data storage (ex: regarding configuration settings), we recommend the use of the file format that [file:consult/1](#) can directly read, that we named, for reference purpose, **ETF** (for *Erlang Term Format*).

This is just a text format for which:

- a line starting with a `%` character is considered to be a comment, and is thus ignored

- other lines are terminated by a dot, and correspond each to an Erlang term (ex: `{base_log_dir, "/var/log"}.}`.)

A basic support for these ETF files is available in

`file_utils:{read,write}_etf_file/*.`

If expecting to read UTF-8 content from such a file, it should:

- have been then opened for writing typically while including the `{encoding,utf8}` option, or have been written with content already properly encoded (maybe more reliable that way)
- start with a `%% -*- coding: utf-8 -*- header`

Support for Code Injection

It may be useful to decide, at compile-time, whether some code should be added / removed / transformed / generated based on **tokens** defined by the user.

This is done here thanks to the use of **conditional primitives** and associated **compilation defines** (sometimes designated as "**macros**", and typically specified in makefiles, with the `-D` flag).

These conditional primitives are gathered in the `cond_utils` module.

As an early example, so that a piece of code prints **Hello!** on the console when executed iff (*if and only if*) the `my_token` compilation token has been defined (through the `-Dmy_token` command-line flag), one may use:

```
cond_utils:if_defined(my_token, io:format("Hello!"))
```

Of course, as such a code injection is done at compilation-time, should compilation defines be modified the modules making use of the corresponding primitives shall be recompiled so that these changes are taken into account.

Let's enter a bit more in the details now.

Defining a token

A *token* (a compilation-time symbol) may or may not be defined.

To define `my_token`, simply ensure that a `-Dmy_token` command-line option is specified to the compiler (ex: refer to `ERLANG_COMPILER_TOKEN_OPT`, in `GNUmakevars.inc`, for an example of definition for these flags).

To define `my_token` and set it to the integer value 127, use the `-Dmy_token=127` command-line option. Values can also be floats (ex: `-Dmy_token=3.14`) or atoms (ex: `-Dmy_token=some_atom`).

A special token is `myriad_debug_mode`; if it is defined at all (and possibly associated to any value), the debug mode of Myriad is enabled.

We recommend that layers built on top of Myriad define their own token for debug mode (ex: `foobar_debug_mode`), to be able to finely select appropriate debug modes (of course all kinds of modes and configuration settings can be considered as well).

Defining the code to inject

Based on the defined tokens, code may be injected; this code can be any Erlang expression, and the value to which it will evaluate (at runtime) can be used as any other value in the program.

Injecting a *single* expression (i.e. not multiple ones) is not a limitation: not only this single expression can be a function call (thus corresponding to arbitrarily many expressions), but more significantly a series of expressions can be nested in a `begin` / `end` block, making them a single expression¹⁰.

¹⁰A previous implementation of `cond_utils` allowed to specify the code to inject either as an expression or as a *list* of expressions. It was actually a mistake, as a single expression to return can be itself a list (ex: `["red", "blue"]`), which bears a different semantics and should not be interpreted as a list of expressions to evaluate. For example, the result from the code to inject may be bound to a variable, in which case we expect `A=["red", "blue"]` rather than `A="red", "blue"` (this latter term being constructed but not used).

Using tokens to enable code injection

Various primitives for *code injection* are available in the `cond_utils` (mostly pseudo-) module¹¹.

There is first `if_debug/1`, to be used as:

```
cond_utils:if_debug(EXPR_IF_IN_DEBUG_MODE)
```

Like in:

```
A = "Joe",
cond_utils:if_debug(io:format("Hello ~s!", [A]))
```

or, to illustrate expression blocks:

```
cond_utils:if_debug(begin
    C=B+1,
    io:format("Goodbye ~p", [C])
end)
```

These constructs will be replaced by the expression they specify for injection, at their location in the program, iff the `myriad_debug_mode` token has been defined, otherwise they will be replaced by nothing at all (hence with exactly *no* runtime penalty; and the result of the evaluation of `if_debug/1` is then not an expression).

Similarly, `if_defined/2`, used as:

```
cond_utils:if_defined(TOKEN, EXPR_IF_DEFINED)
```

will inject `EXPR_IF_DEFINED` if `TOKEN` has been defined (regardless of any value associated to this token), otherwise the `if_defined/2` call will be removed as a whole¹².

As for `if_defined/3`, it supports two expressions:

```
cond_utils:if_defined(TOKEN, EXPR_IF_DEFINED, EXPR_OTHERWISE)
```

For example:

```
% Older versions being less secure:
TLSSupportedVersions = cond_utils:if_defined(us_web_relaxed_security,
    ['tlsv1.3', 'tlsv1.2', 'tlsv1.1', 'tlsv1'],
    ['tlsv1.3'])
```

If `us_web_relaxed_security` has been defined, the first list will be injected, otherwise the second will.

Note that a call to `if_defined/3` results thus in an expression.

Finally, with `if_set_to/{3,4}`, the injection will depend not only of a token being defined or not, but also onto the value (if any) to which it is set.

For `if_set_to/3`:

¹¹Their actual implementation lies in [Myriad's parse transform](#).

¹²So `if_debug(EXPR)` behaves exactly as: `if_defined(myriad_debug_mode, EXPR)`.

```
cond_utils:if_defined(TOKEN, VALUE, EXPR_IF_SET_TO_THIS_VALUE)
```

will inject `EXPR_IF_SET_TO_THIS_VALUE` iff `TOKEN` has been defined and set to `VALUE`. As a result, the specified expression will not be injected if `some_token` has been set to another value, or not been defined at all.

Usage example, `-Dsome_token=42` having possibly been defined beforehand:

```
cond_utils:if_set_to(some_token,42, SomePid ! hello))
```

As for `if_set_to/4`, in:

```
cond_utils:if_set_to(TOKEN, VALUE, EXPR_IF_SET_TO_THIS_VALUE, EXPR_OTHERWISE)
```

`EXPR_IF_SET_TO_THIS_VALUE` will be injected iff `TOKEN` has been defined and set to `VALUE`, otherwise (not set or set to a different value) `EXPR_OTHERWISE` will be.

Example:

```
Level = cond_utils:if_set_to(my_token, foobar_enabled, 1.0, 0.0) + 4.5
```

Finally, the `switch_set_to/{2,3}` primitives allow to generalise these if-like constructs, with one among any number of code branches selected based on the build-time value of a token, possibly with defaults (should the token not be defined at all, or defined to a value that is not among the ones associated to a code branch).

For that we specify a list of pairs, each made of a value and of the corresponding expression to be injected if the actual token matches that value, like in:

```
cond_utils:switch_set_to(TOKEN, [
    {VALUE_1, EXPR_1},
    {VALUE_2, EXPR_2},
    % [...]
    {VALUE_N, EXPR_N}])
```

For example:

```
cond_utils:switch_set_to(my_token, [
    {my_first_value, io:format("Hello!")},
    {my_second_value, begin f(), g(X,debug), h() end},
    {some_third_value, a(X,Y)}])
```

A compilation-time error will be raised if `my_token` is not set, or if it is set to none of the declared values (i.e. not in `[my_first_value, my_second_value, some_third_value]`).

A variation of this primitive exists that applies a default token value if none was, or if the token was set to a value that is not listed among any of the ones designating a code branch, like in:

```
cond_utils:switch_set_to(TOKEN,
    [ {VALUE_1, EXPR_1},
      {VALUE_2, EXPR_2},
      % [...]
      {VALUE_N, EXPR_N}],
    DEFAULT_VALUE)
```

As always with primitives that define a default, alternate branch, they always inject an expression and thus can be considered as such.

For example:

```
ModuleFilename = atom_to_list( cond_utils:switch_set_to(some_token,
    [{1, foo}, {14, bar}, {20, hello}], 14) ++ ".erl"
```

Here, if `some_token` is not defined, or defined to a value that is neither 1, 14 or 20, then the 14 default value applies, and thus `ModuleFilename` is set to `"bar.erl"`.

Refer to [cond_utils_test.erl](#) for further usage examples.

Controlling assertions

It may be convenient that, depending on a compile-time token (ex: in debug mode, typically triggered thanks to the `-Dmyriad_debug_mode` compilation flag), *assertions* (expressions expected to evaluate to the atom `true`) are enabled, whereas they shall be dismissed as a whole should that token not be defined.

To define an assertion enabled in debug mode, use `assert/1`, like in:

```
cond_utils:assert(foo(A,B)==10)
```

Should at runtime the expression specified to `assert/1` be evaluated to a value `V` that is different from the atom `true`, a `{assertion_failed,V}` exception will be thrown.

More generally, an assertion may be enabled by any token (not only `myriad_debug_mode`) being defined, like in:

```
cond_utils:assert(my_token,bar(C))
```

Finally, an assertion may be enabled iff a token (here, `some_token`) has been defined and set to a given value (here, 42), like in:

```
cond_utils:assert(some_token,42,not baz() andalso A)
```

This may be useful for example to control, on a per-theme basis, the level of checking performed, like in:

```
cond_utils:assert(debug_gui,1,basic_testing()),
cond_utils:assert(debug_gui,2,more_involved_testing()),
cond_utils:assert(debug_gui,3,paranoid_testing()),
```

Note that, in this case, a given level of checking should include the one just below it (ex: `more_involved_testing()` should call `basic_testing()`).

Usage Hints

For tokens, at least currently they must be defined as immediate values (atoms); even using a mute variable, like for the `_Default=my_token` expression, or a variable, is not supported (at least yet).

Note that, for primitives that may not inject code at all (ex: `if_debug/1`), if their conditions are not fulfilled, the specified conditional code is dismissed as a whole, it is not even replaced for example by an `ok` atom; this may matter if this conditional is the only expression in a case clause for example, in which case a compilation failure like *"internal error in core; crash reason: function_clause in function v3_core:cxprs/3 called as v3_core:cxprs/..."* will be reported (the compiler sees unexpectedly a clause not having even a single expression).

A related issue may happen when switching conditional flags: it will select/deselect in-code expressions at compile time, and may lead functions and/or variables to become unused, and thus may trigger at least warnings¹³.

For **functions** that could become unused due to the conditional setting of a token, the compiler could certainly be silenced by exporting them; yet a better approach is surely to use:

```
-compile({nowarn_unused_function,my_func/3}).
```

or:

```
-compile({nowarn_unused_function,[my_func/3, my_other_func/0]}).
```

As for **variables**, should A, B or C be reported as unused if `some_token` was not set, then the `basic_utils:ignore_unused/1` function (mostly a no-op) could be of use:

```
[...]
cond_utils:if_defined(some_token,
                      f(A, B, C),
                      basic_utils:ignore_unused([A, B, C])),
[...]
```

Alternatively, `nowarn_unused_vars` could be used instead, at least in some modules.

For more information

Refer for usage and stubs to the `cond_utils` module (defined in [myriad/src/meta](#)), knowing that it is actually implemented thanks to the Myriad parse transform.

For examples and testing, see the `cond_utils_test` module.

¹³Warnings that we prefer promoting to errors, as they constitute a *very* convenient safety net.

Helpers For User Interface Programming

Some services have been defined, in `myriad/src/user-interface`, in order to handle more easily interactions with the user, i.e. to provide a user interface.

Note

The user-interface services, as a whole, are currently *not* functional. A rewriting thereof as been started yet has not completed yet.

Various Flavours of User Interfaces

Such a user interface may be:

- either **text-only**, within a console, relying either on the very basic `text_ui` (for raw text) or its more advanced `term_ui` counterpart (for terminal-based outputs)
- or **graphical**, with `gui`

Text-based user interfaces are quite useful, as they are lightweight, incur few dependencies (if any), and can be used with headless remote servers (`text_ui` and `term_ui` work well through SSH, and require no X server nor mouse).

As for graphical-based user interfaces, they are the richest, most usual, and generally the most convenient, user-friendly interfaces.

The user interfaces provided by Myriad are stateful, they rely on a **state** that can be:

- either **explicit**, in a functional way; thus having to be carried in all calls
- or **implicit**, using - for that very specific need only - the process dictionary (even if we try to stay away of it as much as possible)

We tested the two approaches and preferred the latter (implicit) one, which was found considerably more flexible and thus finally fully superseded the (former) explicit one.

We made our best so that a lower-level API interface (relying on a more basic backend) is **strictly included** in the higher-level ones (ex: `term_ui` adds concepts - like the one of window or box - to the line-based `text_ui`), in order that any program using a given user interface may use any of the next, upper ones as well (provided implicit states are used), in the following order: the `text_ui` API is included in the one of `term_ui`, which is itself included in the one of `gui`.

We also defined the **settings table**, which is a table gathering all the settings specified by the developer, which the current user interface does its best to accommodate.

Thanks to these "Matryoshka" APIs and the settings table, the definition of a more generic `ui` interface has been possible. It selects automatically, based on available local software dependencies, **the most advanced available backend**, with the most relevant settings.

For example a relevant backend will be automatically selected by:

```
$ cd test/user-interface/src
$ make ui_run
```

On the other hand, if wanting to select a specified backend:

```
$ make ui_run CMD_LINE_OPT="--use-ui-backend term-ui"
```

(see the corresponding [GNUmakefile](#) for more information)

Raw Text User Interface: `text_ui`

This is the most basic, line-based monochrome textual interface, directly in raw text with no cursor control.

Located in `{src,test}/user-interface/textual`, see `text_ui.erl` for its implementation, and `text_ui_test.erl` for an example of its use.

Terminal Text User Interface: `term_ui`

This is a more advanced textual interface than the previous one, with colors, dialog boxes, support of locales, etc., based on [dialog](#) (possibly [whiptail](#) could be supported as well). Such backend of course must be available on the execution host then.

For example, to secure these prerequisites:

```
# On Arch Linux:
$ pacman -S dialog
```

```
# On Debian-like distros:
$ apt-get install dialog
```

Located in `{src,test}/user-interface/textual`, see `term_ui.erl` for its implementation, and `term_ui_test.erl` for an example of its use.

Graphical User Interface: `gui`

This interface relied initially on `gs` (now deprecated), now on [wx](#) (a port of [wxWidgets](#)), maybe later in HTML 5 (possibly relying on the [Nitrogen web framework](#) for that). For the base dialogs, [Zenity](#) could have been an option.

Note

GUI services are currently being reworked, to provide a `gs`-like concurrent API while relying underneath on `wx`, with some additions (such as canvases).

The goal is to provide a small, lightweight API (including message types) that are higher-level than `wx`, and do not depend on any particular GUI backend (such as `wx`, `gs`, etc.) to avoid that user programs become obsolete too quickly, as backends for GUI rise and fall relatively often.

So for example the messages received by the user programs shall not mention `wx`, and they should take a form compliant with [WOOPER](#) message conventions, to easily enable user code to rely on WOOPER if wanted.

Located in `{src,test}/user-interface/graphical`, see `gui.erl`,
`gui_color.erl`, `gui_text.erl`, `gui_canvas.erl`, etc., with a few tests
(`gui_test.erl`, `lorenz_test.erl`).

Related information of interest:

- wxErlang: [Getting started](#) and [Speeding up](#), by Arif Ishaq
- <http://wxerlang.dougedmunds.com/>

Myriad's Helper Shell Scripts

A small set of shell scripts has been defined, in `myriad/src/scripts`, to provide generic facilities useful in the context of Myriad¹⁴.

Erlang-Dedicated Scripts

Searching for Erlang elements

ergrep To search for a text pattern through an Erlang source tree.

Usage: `ergrep [-v|--verbose] [-q|--quiet] [-f|--filenames-only] [-i|--insensitive] <Expression to be found in sources> [TARGET_BASE_DIR]`; recursive grep in Erlang source files, either from the TARGET_BASE_DIR directory, if specified, otherwise from the current directory.

Options:

- v or --verbose: be specifically verbose
- q or --quiet: be specifically quiet, just listing matches
- f or --filenames-only: display only filenames, not also the matched patterns, and if there are multiple matches in the same file, its filename will be output only once (implies quiet); useful for scripts
- i or --insensitive: perform case-insensitive searches in the content of files, and also in the searched Erlang filenames

Example: `ergrep -i 'list_to_form(' /tmp`

find-type-definition.sh To search for the definition of a type of interest.

Usage: `find-type-definition.sh [-h|--help] A_TYPE [A_DIR]`; attempts to find the definition of the specified Erlang type from the specified directory (if any), otherwise from the current one.

find-function-specification.sh To search for a function spec of interest.

Usage: `find-function-specification.sh [-h|--help] A_FUNCTION_NAME [A_DIR]`; attempts to find the type specification for the specified Erlang function (even if the name is partial and/or includes an arity) from the specified directory (if any), otherwise from the current one.

find-record-definition.sh To search for the definition of a record of interest.

¹⁴For a more general collection of (different) scripts, one may refer to [Ceylan-Hull](#), notably the [ones to facilitate development](#).

Usage: find-record-definition.sh [-h|--help] A_RECORD_NAME
[A_DIR]: attempts to find the definition of the specified Erlang record
from the specified directory (if any), otherwise from the current one.

Regarding Typing

list-available-types.sh Lists all Erlang types found in specified tree.

Usage: list-available-types.sh [-h|--help] [ROOT_DIR]: lists all types
(according to Erlang type specifications) defined from the ROOT_DIR
directory (if specified) or from the current directory.

add-deduced-type-specs.escript Generates and adds the type specification
of all functions in specified module(s).

Usage: add-deduced-type-specs.escript FS_ELEMENT

Adds, for each selected BEAM file (either specified directly as a file, or
found recursively from a specified directory), in the corresponding source
file(s), for each function, the type specification that could be deduced
from its current implementation.

FS_ELEMENT is either the path of a BEAM file or a directory that
will be scanned recursively for BEAM files.

Note that BEAM files must be already compiled, and with debug
information (see the '+debug_info' compile flag).

Note also that generating type specs this way may not be a good practice.

Regarding Basic Performance Measurement

etop.sh Monitors on the console the busiest Erlang processes of a VM.

Usage: etop.sh [-node NODE_NAME] [-setcookie COOKIE]: shows on
the console the activity of the Erlang processes on specified Erlang node
(enter CTRL-C twice to exit).

Example: etop.sh -node foobar@baz.org -setcookie 'my cookie'

benchmark-command.escript Returns a mean resource consumption for the
specified shell command (one may prefer relying on **benchmark-command.sh**).

Usage: benchmark-command.escript COMMAND: returns a mean
resource consumption for the specified shell command.

Example: benchmark-command.escript "my_script.sh 1 2"

Miscellaneous

make-code-stats.sh Outputs basic statistics about an Erlang code base.

Usage: `make-code-stats.sh [-h|--help] [SOURCE_DIRECTORY]`:
evaluates various simple metrics of the Erlang code found from any
specified root directory, otherwise from the current one.

launch-erl.sh The only shell script on which we rely in order to launch an
Erlang VM.

Usage: `launch-erl.sh [-v] [-c a_cookie] [--sn a_short_node_name | --ln
a_long_node_name | --nn an_ignored_node_name] [--tcp-range
min_port max_port] [--epmd-port new_port] [--fqdn a_fqdn]
[--max-process-count max_count] [--busy-limit kb_size]
[--async-thread-count thread_count] [--background] [--non-interactive]
[--eval an_expression] [--no-auto-start] [-h|--help] [--beam-dir a_path]
[--beam-paths path_1 path_2] [-start-verbatim-options [...]]`: launches
the Erlang interpreter with specified settings.

Detailed options:

- v: be verbose
- c a_cookie: specify a cookie, otherwise no cookie will be
specifically set
- sn a_short_node_name: distributed node using specified short
name (ex: 'my_short_name')
- ln a_long_node_name: distributed node using specified long
name (ex: 'my_long_name')
- nn an_ignored_node_name: non-distributed node, specified
name ignored (useful to just switch the naming options)
- tcp-range min_port max_port: specify a TCP port range for
inter-node communication (useful for firewalling issues)
- epmd-port new_port: specify a specific EPMD port (default:
4369); only relevant if the VM is to be distributed (using short or
long names), initially or at runtime
- fqdn a_fqdn: specify the FQDN to be used
- max-process-count max_count: specify the maximum number of
processes per VM (default: 400000)
- busy-limit size: specify the distribution buffer busy limit, in kB
(default: 1024)
- async-thread-count thread_count: specify the number of
asynchronous threads for driver calls (default: 128)
- background: run the launched interpreter in the background
(ideal to run as a daemon, ex: on a server)
- daemon: run the node as a daemon (relies on `run_erl` and implies
`--background`)

--non-interactive: run the launched interpreter with no shell nor
 input reading (ideal to run through a job manager, ex: on a cluster)
 --eval 'an Erlang expression': start by evaluating this expression
 --no-auto-start: disable the automatic execution at VM start-up
 -h or --help: display this help
 --beam-dir a _path: adds specified directory to the path searched
 for beam files (multiple --beam-dir options can be specified)
 --beam-paths first _path second _path ...: adds specified directories
 to the path searched for beam files (multiple paths can be
 specified; must be the last option)
 --log-dir: specify the directory in which the VM logs (if using
 run _erl) shall be written

Other options will be passed 'as are' to the interpreter with a warning, except if they are listed after a '-start-verbatim-options' option, in which case they will be passed with no warning.

If neither '--sn' nor '--ln' is specified, then the node will not be a distributed one.

Example: launch-erl.sh -v --ln ceylan --eval
 'class _TimeManager _test:run()'

show-xml-file.erscript Displays the content of the specified XML file.

Usage: show_xml_file.erscript XML_FILE_PATH
 Displays sequentially in a {name,Value} tree the structure of specified
 XML file (XML elements along with their XML attributes).

More General Scripts

To generate documentation These scripts are mostly unrelated to Erlang, yet are useful to be available from our most basic layer (Myriad).

generate-docutils.sh Generates a proper PDF and/or HTML file from specified RST (**reStructuredText**) one (main, standalone script).

Usage: generate-docutils.sh <target rst file>
 [--pdf|--all]<comma-separated path(s) to CSS file to be used, ex:
 common/css/XXX.css,other.css>] [--icon-file ICON_FILENAME]

Generates a final document from specified docutils source file (***.rst**).

If '--pdf' is specified, a PDF will be created, if '--all' is specified, all output formats (i.e. HTML and PDF) will be created, otherwise HTML files only will be generated, using any specified CSS file.

generate-pdf-from-rst.sh Generates a proper PDF and/or HTML file from specified RST ([reStructuredText](#)) one; the previous **generate-docutils.sh** script is often preferred to this one, which depends on Myriad.

Usage: generate-pdf-from-rst.sh RST_FILE: generates a PDF file from the specified RST file, overwriting any past file with that name.

Ex: 'generate-pdf-from-rst.sh my_file.rst' will attempt to generate a new 'my_file.pdf' file.

Script-based Apps

These shell scripts are actually user-facing shell interfaces that plug directly on some more involved Erlang programs, i.e. applications that are [available here](#).

generate-password.sh Generates a proper random password respecting various rules, whose simple application can be transparently checked (probably at least more easily audited than most password managers - thus maybe more trustable).

Usage: generate-password.escript [-a ALPHABET|--alphabet ALPHABET]
[-l MIN_LEN MAX_LEN|--length MIN_LEN MAX_LEN]
[-h|--help]

Generates a suitable password, where:

- ALPHABET designates the set of characters to draw from (default one being 'extended'), among:

- * 'base': alphanumeric letters, all cases [A-Za-z0-9]
- * 'extended': 'base' + basic punctuation (i.e. '[](){}:;,-_!?.'))
- * 'full': 'base' + all punctuation (i.e. basic + "'@ /&\$%^%+=+|'")

- MIN_LEN and MAX_LEN are the respective minimum and maximum numbers of characters

(bounds included) used to generate this password [default: between 15 and 20]

See also: the [security section](#) of Ceylan-Hull, for more general guidelines and tooling regarding the proper management of credentials.

merge.sh Helps merging efficiently and reliably file trees; it is actually a rather involved text-based application that allows scanning/comparing/merging

trees, typically in order to deduplicate file hierarchies that were cloned once yet may have since then diverged.

Usage: following operations can be triggered:

- 'merge.sh --input INPUT_TREE --reference REFERENCE_TREE'
- 'merge.sh --scan A_TREE'
- 'merge.sh --rescan A_TREE'
- 'merge.sh --resync A_TREE'
- 'merge.sh --uniquify A_TREE'
- 'merge.sh -h' or 'merge.sh --help'

Ensures, for the first form, that all the changes in a possibly more up-to-date, "newer" tree (INPUT_TREE) are merged back to the reference tree (REFERENCE_TREE), from which the first tree may have derived. Once executed, only a refreshed, complemented reference tree will exist, as the input tree will have been removed: all its original content (i.e. its content that was not already in the reference tree) will have been transferred in the reference tree.

In the reference tree, in-tree duplicated content will be either kept as it is, or removed as a whole (to keep only one copy thereof), or replaced by symbolic links in order to keep only a single reference version of each actual content.

At the root of the reference tree, a '.merge-tree.cache' file will be stored, in order to avoid any later recomputations of the checksums of the files that it contains, should they have not changed. As a result, once a merge is done, the reference tree may contain an uniquified version of the union of the two specified trees, and the input tree will not exist anymore after the merge.

For the second form (--scan option), the specified tree will simply be inspected for duplicates, and a corresponding '.merge-tree.cache' file will be created at its root (to be potentially reused by a later operation).

For the third form (--rescan option), an attempt to rebuild an updated '.merge-tree.cache' file will be performed, computing only the checksum of the files that were not already referenced, or whose timestamp or size changed.

For the fourth form (--resync option), a rebuild even lighter than the previous rescan of '.merge-tree.cache' will be done,

checking only sizes (not timestamps), and updating these timestamps.

For the fifth form (`--uniquify` option), the specified tree will be scanned first (see the corresponding operation), and

For the fifth form (`--uniquify` option), the specified tree will be scanned first (see the corresponding operation), and then the user will be offered various actions regarding found duplicates (being kept as are, or removed, or replaced with symbolic links), and once done a corresponding up-to-date `'.merge-tree.cache'` file will be created at its root (to be potentially reused by a later operation).

For the sixth form (`-h` or `--help` option), displays this help.

Note that the `--base-dir A_BASE_DIR` option can be specified by the user to designate the base directory of all relative paths mentioned. When a cache file is found, it can be either ignored (and thus recreated) or re-used, either as it is or after a weak check, where only file existence, sizes and timestamps are then verified (not checksums).

See also: the `test-all` target of the merge-related [makefile](#), to give it a try before applying such procedure to your data of interest.

Utility Toolbox

This is the **core** of the **Ceylan-Myriad** library: a toolbox comprising many helper functions (with their tests), defined in the `myriad/src/Utils` directory, often providing enhanced, more specialised services compared to the ones offered by the Erlang standard libraries.

These helpers (code and typing information) are thematically aggregated in modules that are generally suffixed by `_utils`, and include:

- many **basic, general-purpose services**, defined in `basic_utils.erl`, regarding:
 - the base types we defined
 - notifications
 - message handling
 - many miscellaneous functions
- **cipher**-related facilities (basic, a bit exotic chained symmetric encryptions, notably with Mealy machines), in `cipher_utils.erl`
- functions to manage Erlang **compiled BEAM code** (`code_utils.erl`)
- services to manage the **execution of other programs** (`executable_utils.erl`), to:
 - locate said executables
 - to execute functional services (ex: display a PDF) regardless of the actual executable involved
- a few **test-related facilities**, in `test_facilities.erl`
- services to handle more easily the (UNIX) shells and also the command-line arguments (a bit like `getopt`), regardless of the interpreter or escript context (`shell_utils.erl`)
- helpers for **file-based** I/O operations (`file_utils.erl`); note that we now recommend not to open files with a specific encoding being set, but instead to encode the content before any writing thereof; refer to the **Regarding encodings and Unicode** section in `file_utils.erl` for further information
- services to manage identifiers of various sorts, including sortable ones (`id_utils.erl`)
- a very basic support of **Finite State Machines** (`fsm_utils.{e,h}.erl`)
- a few operations defined on **graphs** (`graph_utils.erl`, with `find_breadth_first/3,4`)
- extra operations defined on **lists** (`list_utils.erl`), including rings
- support for **network**-related concerns (`net_utils.erl.{e,h}.erl`)

- services to offer **randomness** (`random_utils.erl`), with regard to various sources (the Erlang built-in algorithm, `crypto`, newer ones like `exsplus` - our current default, `exs64` and `exs1024`), for seeding, drawing, etc.
- very little support of **RDF** operations, standing for [Resource Description Framework](#) (`rdf_utils.erl`)
- facilities to handle content to the web, to HTTP, etc. (`web_utils.erl`) and to perform **REST calls** (`rest_utils.erl`), using built-in `httpc` and `http_client`, including JSON services
- elements for the sending of **SMS** (`sms_utils.erl`), based either on third-party providers providing REST APIs, or via a mobile phone (typically connected thanks to a USB cable); nothing as advanced as [Ceylan-Mobile](#), though
- support for operations at the **operating-system** level (`system_utils.{e,h}.erl`)
- services to handle **text** (`text_utils.erl`)
- services to handle **binary information**, such as CRC (`bin_utils.erl`)
- functions to manage **time** information (`time_utils.erl`)
- a few helpers to ease the writing of [escripts](#) relying on the Myriad layer (`script_utils.erl`)
- services addressed to the use of OTP, in `otp_utils.erl`, allowing notably to run OTP applications out of an OTP context (typically through our native build/run system rather than through `rebar3` and even OTP releases)
- services about all kinds of **units** (`unit_utils.erl`); refer to the [Management of Units](#) section below for more information
- support for **CSV** (Comma-Separated Values) files (`csv_utils.erl`)
- support for **JSON** documents (`json_utils.erl`), based on any available parser backend (see the [Myriad-level Third-Party Dependencies](#) section for more details)
- basic services for **trace emission** (a.k.a logging - not related to Erlang tracing), either directly through `trace_utils.erl`, or thanks to `trace_bridge.erl` - typically to rely on more advanced trace systems such as [Ceylan-Traces](#); now compliant with newer OTP logger and Syslog protocol as defined in [RFC 5424](#), collecting both userland traces and VM-level logs
- very basic facilities for **applications** (not in the sense of OTP ones), in `app_facilities.{e,h}.erl` with an example (`most_basic_example_app.erl`)
- a bit of **locale management**, in `locale_utils.erl`
- minor services about the **monitoring of Erlang processes**, in `monitor_utils.erl` and their **registering** in naming services, in `naming_utils.erl`

- facilities to better **interface Erlang to other languages**, in `language_utils.erl` and `{python,java}_utils.erl`; nothing as advanced as [Ceylan-Seaplus](#), though

Support for Metaprogramming

Over time, quite a lot of developments grew to form primitives that manage ASTs (*Abstract Syntax Trees*), based on Erlang's parse transforms. These developments are gathered in the `src/meta` directory, providing notably:

- `meta_utils.{e,h}rl`: basic primitives to **transform ASTs**, with a bit of testing (`meta_utils_test`)
- `type_utils`: a still rather basic toolbox to **manage data types** - whether built-in, compound or parametrised (expressed as strings, as terms, etc.)
- `ast_*` modules to handle the various **elements that can be found in an AST** (ex: `ast_expression`, `ast_type`, `ast_pattern`, etc.)

Finally, a few usage examples of these facilities are:

- `minimal_parse_transform_test`: the simplest parse transform test that we use, typically operating on `simple_parse_transform_target`
- `example_parse_transform`: a rather minimal parse transform
- `myriad_parse_transform`: the parse transform used within Myriad, transforming each and every module of that layer (and of at least some modules of upper layers)

So the purpose of this parse transform is to **convert ASTs that are Myriad-compliant into ASTs that are directly Erlang compliant**.

For that, following changes are operated:

- in type specifications, the Myriad-specific `void/0`, `maybe/1` and `fallible/{1,2}` types are adequately translated:
 - `void()` becomes `basic_utils:void()`, a type alias of `any()`, made to denote returned terms that are not expected to be used by the caller (as if that function's only purpose was its side-effects)
 - `maybe(T)` becomes the type union `'undefined'|T`
 - `fallible(T)` becomes ultimately the type union `{'ok',T}|{'error',term()}`, while `fallible(Tok, TError)` becomes `{'ok',Tok}|{'error',TError}`
- both in type specifications and actual code, `table/2`, the Myriad-specific associative table pseudo-type, is translated into an actual [table type](#)
- the `const_table` module allows to generate sharable, read-only associative tables probably in the most efficient way in Erlang (refer to the [const table](#) section for further information)
- the `cond_utils` services drive conditional code injection: based on the build-time tokens defined, their values can be used to perform compilation-time operations such as `if` (see in this module `if_debug/1`, `if_defined/{2,3}`, `if_set_to/{3,4}`), `switch` (see `switch_set_to/{2,3}`, possibly with a default clause) or `assert` (`assert/{1,2,3}`); if useful, it should be fairly easy (infrastructure mostly ready) to transform the (currently constant)

user-defined build tokens into mutable variables and to add for example compile-time assignments (`cond_utils:create_token(TOKEN, MAYBE_INITIAL_VALUE)`, `cond_utils:set_token_value(TOKEN, VALUE)` and `cond_utils:remove_token(TOKEN, MAYBE_INITIAL_VALUE)`) of these variables and loops (**for**, **while**, etc.) if not going for a Turing-complete language, if ever that made sense for some uses; see the [Support for Code Injection](#) for additional usage details regarding the supported primitives

More generally, Myriad offers the support to traverse *any* AST (the whole Erlang grammar is supported, in its abstract form) and to **transform** it (ex: an expression being removed, transformed or replaced by other expressions), with the ability for the user to define his own type/call replacement mappings, or more general transformation functions to be triggered when specified elements are found in the AST (ex: remote calls with relevant MFA).

The traversal may be done in a stateful manner, i.e. any user-defined transformation will be able to access (read/write) any state of its own in the course of the traversal.

As a result, a single pass through the input AST may be done, in which any kind of transformations may be applied, resulting in another (Erlang-compliant) AST being output and afterwards compiled.

Management of Units

Motivation

A value of a given type (ex: a float) can actually correspond to quantities as different as meters and kilowatts per hour.

Therefore **units shall preferably be specified alongside with values being processed**, and a language to express, check and convert these units must be retained. Of course units are of interest as other metadata are - such as accuracy, semantics, etc.

Available Support

The Myriad layer provides such a service, in a very basic, ad hoc form (which is useful to introduce "special" non-physical, non-standard units, such as `euro/year`), meant to be enriched over time.

Specifying Units

Aliases For convenience, *aliases* of units can be defined, i.e. alternate names for a given canonical unit. For example the Hertz unit (Hz) is an alias of the s^{-1} (per-second) canonical unit.

Built-in Units So one may use the following **built-in units**, whose symbol¹⁵ is specified here between brackets, like in "`N.m`" (an alternate notation is to prefix a unit with `U:`, like in "`U: N.m`"):

- the seven [SI base units](#), namely:
 - meter, for length [m]
 - gram, for mass [g]¹⁶ (note: this is a footnote, not an exponent!)
 - second, for time [s]
 - ampere, for electric current [A]
 - kelvin, for thermodynamic temperature [K]
 - mole, for the amount of substance [mol]
 - candela, for luminous intensity [cd]
- the usual [derived units](#), notably:
 - hertz, for frequency [Hz]
 - degree, for degree of arc [°] (not supported yet)
 - radian, for angle [rad] (not supported yet)
 - steradian, for solid angle [sr] (not supported yet)
 - newton, for force, weight [N]
 - pascal, for pressure, stress [Pa]
 - joule, for energy, work, heat [J]
 - watt, for power, radiant flux [W]

- coulomb, for electric charge, quantity of electricity [C]
 - volt, for voltage, electrical potential difference, electromotive force [V]
 - farad, for electrical capacitance [F]
 - ohm, for electrical resistance, impedance, reactance [Ohm]
 - siemens, for electrical conductance [S]
 - weber, for magnetic flux [Wb]
 - tesla, for magnetic field strength, magnetic flux density [T]
 - henry, for inductance [H]
 - lumen, for luminous flux [lm]
 - lux, for illuminance [lx]
 - becquerel, for radioactive decays per unit time [Bq]
 - gray, for absorbed dose of ionizing radiation [Gy]
 - sievert, for equivalent dose of ionizing radiation [Sv]
 - katal, for catalytic activity [kat]
- the units **widely used** in conjunction with SI units (note that they may not respect the principle of being a product of integer powers of one or more of the base units):
 - litre, for 10^{-3}m^3 volumes [L]
 - tonne, for 1,000 kilogram masses [t]
 - electronvolt, for $1.602176565(35) \cdot 10^{-19}$ joule energies [eV]
 - minute, for 60-second durations [min]
 - hour, for 60-minute durations [h]
 - day, for 24-hour durations [day]
 - week, for 7-day durations [week]
 - the **special** units (they generally cannot map directly to any SI unit, yet can be handled separately), designating:
 - month [month] (correspondence to base time units unspecified, as this duration is not constant; ex: a month can be 29, 30 or 31 days)
 - year [year] (correspondence to base time units unspecified, as this duration is not constant; ex: a year can be 365, 366 or 365.25 days, etc.)
 - degree Celsius, for temperature relative to 273.15 K [°C] (see note below)
 - dimension-less quantities (ex: an index) [dimensionless] (most probably clearer than m/m)
 - a count, i.e. a dimensionless number, generally a positive integer [count] (ex: 14), considered as an alias of dimensionless

- a ratio, i.e. a dimensionless floating-point value, generally displayed as a percentage [**ratio**] (ex: -12.9%); another alias of **dimensionless**
- currencies, either [**\$**] (US Dollar) or [**euros**] (Euro), whose exchange rates of course vary
- values whose unit has not been specified [**unspecified_unit**]
- **metric prefixes** thereof, i.e. multiples and sub-multiples of the units previously mentioned; currently the supported prefixes are:
 - yotta, i.e. 10^{24} [**Y**]
 - zetta, i.e. 10^{21} [**Z**]
 - exa, i.e. 10^{18} [**E**]
 - peta, i.e. 10^{15} [**P**]
 - tera, i.e. 10^{12} [**T**]
 - giga, i.e. 10^9 [**G**]
 - mega, i.e. 10^6 [**M**]
 - kilo, i.e. 10^3 [**k**]
 - hecto, i.e. 10^2 [**h**]
 - deca, i.e. 10 [**da**]
 - deci, i.e. 10^{-1} [**d**]
 - centi, i.e. 10^{-2} [**c**]
 - milli, i.e. 10^{-3} [**m**]
 - micro, i.e. 10^{-6} [**µ**]
 - nano, i.e. 10^{-9} [**n**]
 - pico, i.e. 10^{-12} [**p**]
 - femto, i.e. 10^{-15} [**f**]
 - atto, i.e. 10^{-18} [**a**]
 - zepto, i.e. 10^{-21} [**z**]
 - yocto, i.e. 10^{-24} [**y**]

Note

There is a problem with temperatures, as they can be expressed at least in kelvins or degrees Celsius, whereas the two corresponding scales do not match, since there is an offset: $[K] = [^{\circ}C] + 273.15$.

As a result, unit conversions would require updating as well the corresponding value, and, more generally, they should be treated as fully distinct units (ex: $kW/^{\circ}C$ cannot be automatically converted in terms of SI base units, i.e. using K).

This is why we "degraded" Celsius degrees, from a derived unit to a special one.

The same applies to the Fahrenheit unit (a likely addition), as: $[^{\circ}C] = 5/9 \cdot ([^{\circ}F] - 32)$.

¹⁵To avoid requesting the user to type specific Unicode characters, we transliterated some of the symbols. For example, instead of using the capital Omega letter, we used **Ohm**.

¹⁶We preferred here deviating a bit from the SI system, by using this non-prefixed unit (the *gram*) instead of the SI standard one, which happens to be the *kilogram*.

Composing One's Units So an actual unit can be composed from the aforementioned built-in units (be they base, derived, widely used, special units; prefixed or not)¹⁷, using two built-in operators, which are "." (multiply, represented by the dot character - not "*") and "/" (divide, represented by the forward slash character).

The resulting type shall be specified as a string, containing a series of built-in units (potentially prefixed) alternating with built-in operators, like in:
`"kW.s/m".`

Note

As a result, `"kWh"` is not a valid unit: it should be denoted as `"kW.h"`. Similarly, `"W/(m.k)"` is not valid, since parentheses are currently not supported: `"W/m/k"` may be used instead.

Finally, exponents can be used as a shorthand for both operators (ex: `kg.m^2.s^-1`, instead of `kg.m.m/s`). They should be specified explicitly, thanks to the caret character ("`^`"); for example `"m^2/s"`, not `"m^2/s"`. If deemed both safe and useful, we may consider in the future performing:

- symbolic unit checking (i.e. determining that a derived unit such as `N.s` (newton.second) is actually, in canonical SI units, `m^2.kg.s^-1`), and thus that values of these two types can safely be used indifferently in computations
- automatic value conversions (ex: converting `km/hour` into `m/s`), provided that the overall computational precision is not significantly deteriorated

The corresponding mechanisms (type information, conversion functions, unit checking and transformation, etc.) are defined in `unit_utils.erl` and tested in `unit_utils_test.erl`, in the `myriad/src/units` directory.

Checking Units A typical example:

```
1> MyInputValue="-24 mS.m^-1".
2> {Value,Unit}=unit_utils:parse_value_with_unit(MyInputValue).
3> io:format("Corresponding value: ~f.~n", [ Value ] ).
Corresponding value: -24.0.
4> io:format("Corresponding unit: ~s.~n",
    [unit_utils:unit_to_string(Unit)]).
"s^3.A^2.g^-1.m^-3, of order -6"
5> unit_utils:value_with_unit_to_string(Value,Unit).
"-2.4e-5 s^3.A^2.g^-1.m^-3"
```

Possible Improvements Regarding Dimensional Analysis

Some programming languages provide systems to manage dimensional information (ex: for physical quantities), generally through add-ons or libraries (rarely as a built-in feature).

A first level of support is to provide, like here, an API to manage units. Other levels can be:

¹⁷In the future, defining an actual unit from other actual units might be contemplated.

1. to integrate unit management directly, seamlessly in language expressions, as if it was built-in (as opposed to having to use explicitly a third-party API for that); for example at least half a dozen different libraries provide that in Python
2. to be able to define "polymorphic units and functions", for example to declare in general that a speed is a distance divided by a duration, regardless of the possible units used for that
3. to perform *static* dimensional analysis, instead of checking units at runtime

The two latter use cases can for example be at least partially covered by Haskell libraries.

SQL support

About SQL

Some amount of [SQL](#) (*Structured Query Language*) support for relational database operations is provided by the Myriad layer. As this support relies on an optional prerequisite, this service is disabled by default.

Database Back-end

To perform SQL operations, a corresponding software solution must be available.

The SQL back-end chosen here is the [SQLite 3](#) library. It provides a self-contained, serverless, zero-configuration, transactional SQL database. It is an embedded SQL database engine, as opposed to server-based ones, like [PostgreSQL](#) or [MariaDB](#).

It can be installed on Debian thanks to the `sqlite3` and `sqlite3-dev` packages, `sqlite` on Arch Linux..

We require version 3.6.1 or higher (preferably: latest stable one). It can be checked thanks to `sqlite3 --version`.

Various related tools are very convenient in order to interact with a SQLite database, including `sqlitebrowser` and `sqlliteman`.

On Arch Linux, one can thus use: `pacman -Sy sqlite sqlitebrowser sqlliteman`.

Testing the back-end:

```
$ sqlite3 my_test

SQLite version 3.13.0 2016-05-18 10:57:30
Enter ".help" for usage hints.
sqlite> create table tblone(one varchar(10), two smallint);
sqlite> insert into tblone values('helloworld',20);
sqlite> insert into tblone values('my_myriad', 30);
sqlite> select * from tblone;
helloworld|20
my_myriad|30
sqlite> .quit
```

A file `my_test`, identified as SQLite 3.x database, must have been created, and can be safely removed.

Erlang SQL Binding

This database system is directly accessed thanks to an Erlang binding.

Two of them have been identified as good candidates:

- [erlang-sqlite3](#): seems popular, with many contributors and users, actively maintained, based on a `gen_server` interacting with a C-node, involving only a few source files
- [esqlite](#): based on a NIF, so more able to jeopardize the stability of the VM, yet potentially more efficient

Both are free software.

We finally preferred `erlang-sqlite3`.

By default we consider that this back-end has been installed in
~/Software/erlang-sqlite3. The `SQLITE3_BASE` variable in
myriad/GNUMakevars.inc can be set to match any other install path.

Recommended installation process:

```
$ mkdir ~/Software
$ cd ~/Software
$ git clone https://github.com/alexeyr/erlang-sqlite3.git
Cloning into 'erlang-sqlite3'...
remote: Counting objects: 1786, done.
remote: Total 1786 (delta 0), reused 0 (delta 0), pack-reused 1786
Receiving objects: 100% (1786/1786), 3.24 MiB | 570.00 KiB/s, done.
Resolving deltas: 100% (865/865), done.
Checking connectivity... done.
$ cd erlang-sqlite3/
$ make
rm -rf deps ebin priv/*.so doc/* .eunit/* c_src/*.o config.tmp
rm -f config.tmp
echo "normal" > config.tmp
./rebar get-deps compile
==> erlang-sqlite3 (get-deps)
==> erlang-sqlite3 (compile)
Compiled src/sqlite3_lib.erl
Compiled src/sqlite3.erl
Compiling c_src/sqlite3_drv.c
[...]
```

Testing the binding:

```
make test
./rebar get-deps compile eunit
==> erlang-sqlite3 (get-deps)
==> erlang-sqlite3 (compile)
==> erlang-sqlite3 (eunit)
Compiled src/sqlite3.erl
Compiled src/sqlite3_lib.erl
Compiled test/sqlite3_test.erl
===== EUnit =====
module 'sqlite3_test'
  sqlite3_test: all_test_ (basic_functionality)...[0.002 s] ok
  sqlite3_test: all_test_ (table_info)...ok
  [...]
  sqlite3_lib: delete_sql_test...ok
  sqlite3_lib: drop_table_sql_test...ok
  [done in 0.024 s]
  module 'sqlite3'
=====
All 30 tests passed.
Cover analysis: ~/Software/erlang-sqlite3/.eunit/index.html
```

Pretty reassuring.

SQL Support Provided By the *Myriad* Layer

To enable this support, once the corresponding back-end (see [Database Back-end](#)) and binding (see [Erlang SQL Binding](#)) have been installed, the `USE_SQLITE` variable should be set to `true` in `myriad/GNUMakevars.inc` and Myriad shall be rebuilt.

Then the corresponding implementation (`sql_support.erl`) and test (`sql_support_test.erl`), both in `myriad/src/data-management`, will be built (use `make clean all` from the root of Myriad) and able to be run (execute `make sql_support_run` for that).

Testing it:

```
$ cd myriad/src/data-management
$ make sql_support_run
    Compiling module sql_support.erl
    Compiling module sql_support_test.erl
    Running unitary test sql_support_run
[...]
--> Testing module sql_support_test.
Starting SQL support (based on SQLite3).
[...]
Closing database.
Stopping SQL support.
--> Successful end of test.
(test finished, interpreter halted)
```

Looks good.

SQL-related Troubleshooting

Compiling module `sql_support.erl` : can't find include file `"sqlite3.hrl"`

- `USE_SQLITE` not set to `true` in `myriad/GNUMakevars.inc`
- `erlang-sqlite3` back-end not correctly installed (ex: `SQLITE3_BASE` not pointing to a right path in `myriad/GNUMakevars.inc`)

Myriad Main Conventions

We list here the conventions of all sorts that the Myriad code (base or contributed one) - and also code in the software stack based on it - shall obey.

Text Conventions

The purpose here is to ensure a sufficient **code homogeneity**; for example in all source files are in such a "canonical form", analysing their differences (`diff`) is made simpler.

Any text editor can be used, provided that it saves source files with the UNIX, not DOS, conventions (i.e. lines terminating by the LF character, not by the CRLF characters).

The use of syntax highlighting is encouraged.

Recommended text editors are:

- [Emacs](#) (see our [init.el](#))
- [Visual Studio Code](#) (a.k.a. Vscodex)
- [ErlIDE](#) (based on Eclipse)
- Vim, IntelliJ, Gedit, Nedit, etc.

The main editors integrate the *Language Server Protocol* (also known as LSP), so [Erlang LS](#) can be used with them ([Flycheck](#) can be another option). For that we recommend the use of our [erlang_ls.config](#) file.

Source files should be formatted for a 80-character width: no character should be present after the 79th column of a line.

Except in very specific cases, only ASCII code should be used (ex: no accented characters).

Tabulations should be preferred to series of spaces, and the text should be formatted according to 4-character tabulations.

All redundant whitespaces should be removed, preferably automatically (see the Emacs `whitespace-cleanup` command). This is why, with the [emacs settings](#) that we recommend, pressing the F8 key removes for example the yellow areas in the current buffer by replacing any series of four spaces by a corresponding tabulation.

We would prefer that all files (especially source ones; including the contributed ones) are "whitespace-clean" before being committed. As mentioned, such a transformation can be done directly from Emacs. If using another editor, please ensure that the [fix-whitespaces.sh](#) script has been run on the target sources (possibly automatically thanks to a VCS hook) *before* committing them; the [fix-whitespaces-in-tree.sh](#) script may be also used, in order to perform a bulk transformation.

All elements of documentation should be written in English, possibly translated to other languages. Spell-checking is recommended.

Coding Practices

In terms of coding style, we would like that the sources remain as uniform as possible, regarding naming, spacing, code/comments/blank line ratios.

We would like that, roughly and on average, the same ratio applies for blank lines, comments and lines of code.
For that one may use the either directly `make stats` from the root of the layer or the [make-code-stats.sh](#) script.
For example:

```
In the Erlang source code found from XXX/ceylan/myriad, we have:
+ 160 source files (*.erl), 54 header files (*.hrl)
+ a grand total of 89151 lines:
- 27186 of which (30.4%) are blank lines
- 28751 of which (32.2%) are comments
- 33214 of which (37.2%) are code
```

The most obvious conventions are:

- the **settings of the build chain** should be used (ex: with regard to compiler flags) and adapted/completed if needed; the (possibly-specialised) `GNUmakesettings.inc`, `GNUmakerules.inc` and `GNUmakevars.inc` files should be relied upon
- **no warning should be tolerated**; anyway our build chain treats warnings as (blocking) errors
- **test cases** should be developed alongside most if not all modules; ex: if developing `src/class_X.erl`, then probably the `test/class_X_test.erl` testing code should be developed, after or, even preferably, before the implementation of the tested code; test success should be evaluated automatically, by the code (ex: thanks to pattern matching), not by the person running the test (ex: who would have to compare visually the actual results with the expected ones); in some cases, only **integrated tests** can be devised in practice; tests should be gathered in **test suites**, that should be runnable automatically (`make test`, recursively through child directories) and fail loudly (and in a blocking manner) at the first error met
- **multiple levels of quality documentation** should be made available to the code user, and probably be written in parallel to the code; there are at least three documentation levels:
 - lower-level documentation: code should always be **densely commented**, with documentation headers added to all functions, inlined comments (not paraphrasing the code) and self-describing symbols: function names, variable names (ex: `RegisteredState = ...` to be preferred to `NewState = ...`), etc.; more generally all names shall be long enough to be descriptive (clarity preferred over compactness); type specifications also pertain to this low-level documentation effort
 - higher-level **design and/or implementation notes**: they should be available as a set of paragraphs in each source file, before the function definitions, to describe the context and constraints, and help understanding how the features are implemented, and why

- high-level **developer and user documentation** should be made available, targeting at least HTML and PDF outputs, possibly offering a wiki access as well
- more generally, **comments** should be clear and precise, numerous, rich and complete (overall, in terms of line counts, we target roughly 1/3 of code, 1/3 of blank lines and 1/3 of comments); all comments shall be written in UK English, start with a single % and be properly word-wrapped (use `meta-q` with our Emacs settings to take care of that automatically)
- **indentation** should respect, as already explained, the 80-character width and 4-space tabulation; however the default built-in Erlang indentation mode of `emacs` can hardly be used for that, as it leads to huge width offsets (we may in the future provide a variation of the `elisp` code for the emacs indentation rules)
- **spacing homogeneity** across source files should be enforced; for example three blank lines should exist between two function definitions, one between the clauses of any given function (possibly two in case of longer clauses); for the sake of "visual parsing", arguments should be separated by spaces (ex: `f(X) -> ...`, not `f(X) -> ...`), especially if they are a bit complex (`f(A={U,V}, B, _C) -> ...`, not `f(A={U,V},B,_C) -> ...` or the usual `f(A={U,V}, B, _C) -> ...`)
- for **type-related conventions**, at least all exported functions shall have a `-spec` declaration; if an actual type is referenced more than once (notably in a given module), a specific user-defined type shall be defined; types shall be defined in "semantic" terms rather than on technical ones (ex: `-type temperature() :: ...` than `float()`; developers may refer to, or enrich, `myriad/src/utils/unit_utils.erl` for that)
- the **latest stable version of Erlang** should be used, preferably built thanks to our `myriad/conf/install-erlang.sh` script
- the official *Programming Rules and Conventions* should be enforced, as defined [here](#)
- the function definitions shall follow **the same order** as the one of their exports
- helper functions **shall preferably be identified as such**, with an `(helper)` comment
- if an helper function is specific to an exported function, it shall be defined just after this function; otherwise it should be defined in the **helper section**, placed just after the definition of the exported functions
- defining distinct (non-overlapping), explicit (with a clear-enough name), numerous (statically-defined) **atoms** is cheap; each atom found in the sources is generally to be involved in at least one type definition
- the use of `case ... of ... end` should be preferred to the use of `if` (never used in our code base)

- we also prefer that the various patterns of a case are indented with exactly one tabulation, and that the closing **end** lies as much as possible on the left (ex: if having specified `MyVar = case ... end`, then **end** should begin at the same column as `MyVar`); the same applies to `try ... catch ... end` clauses
- when a term is ignored, instead of using simply `_`, one should define a **named mute variable** in order to provide more information about this term (ex: `_TimeManagerPid`); one should then to accidental matching of such names (now a warning is emitted)
- some conventional variable names are, and may be, extensively used: **Res** for result, **H** and **T** for respectively the head and tail of a list on which we recursively iterate
- when needing an **associative table**, use the **table** pseudo-module; a key/value pair shall be designated as a table *entry* (ex: variable named as `RoadEntry`)
- regarding the in-code management of **text**:
 - if a text is to be rather static (constant) and/or if it is to be exchanged between processes, then it should be a UTF8 **binary**, and its type shall be declared as `text_utils:bin_string()`
 - other, a plain string (`text_utils:usttring()`) shall be used
- when defining a non-trivial datastructure, a **record** shall be used (rather than, say, a mere ad-hoc tuple or a map of undocumented structure...), a corresponding **type** should be then defined (ex: a `foobar` record leading to a `foobar()` type), and a **function to describe it** as text shall be provided (ex: `-spec foobar_to_string(foobar()) -> text_utils:string()`)
 - **mute variables** should be used as well to document actual parameters; for example `f(3,7,10)` could preferably be written as a clearer `f(_Min=3,_Max=7,_Deviation=10)`

Note

Mute variables are however actually bound, thus if for example there is in the same scope `_Min=3` and later `_Min=4`, then a badmatch will be triggered at runtime; therefore names of mute variables should be generally kept unique in a given scope.

- type shorthands may be defined; for example, if using repeatedly within a module `text_utils:usttring()`, a local, non-exported type shorthand (`-type ustring() :: text_utils:usttring()`) may be defined so that all other uses of this type become simply `ustring()`

As not all typos may be detected at compilation-time (ex: wrong spelling for a module), we recommend, for source code, the use of additional static checkers, as discussed in the [type-checking](#) section.

Execution Targets

Two execution target modes have been defined:

- **development** (the default): meant to simplify the task of developers and maintainers by reporting as much information and context as possible, even at the expense of some performances and reliability (ex: no retry in case of failure, shorter time-outs not to wait too long in case of errors, more checks, etc.)
- **production**: mostly the reciprocal of the **development** mode, whose purpose is to favor efficient, bullet-proof operations

These execution targets are *compile-time* modes, i.e. they are set once for all when building the layer at hand (probably based, if using OTP, on the rebar corresponding modes - respectively **dev** and **prod**).


See `EXECUTION_TARGET` in `GNUmakevars.inc` to read and/or set them.

The current execution target is of course available at runtime on a per-layer level, see `basic_utils:get_execution_target/0` for more information.

This function shall be compiled once per layer to be accurate, in one of its modules. It is just a matter of adding the following include in such module:

```
-include_lib("myriad/utils/basic_utils.hrl").
```

Geometric Conventions



For **space** coordinates, three axes are defined for a global referential:

- abscissa: X axis (in red, #FF0000)
- ordinate: Y axis (in green, #008000)
- depth: Z axis (in blue, #0000FF)

For each of them, generally 1.0 corresponds to 1 meter, otherwise to 1 [light-second](#) (i.e. roughly 300 000 km¹⁸).

¹⁸Then for more human-sized distances, a scale of one light-nanosecond (10^{-9} second) might be more convenient, as it corresponds almost to 30 cm.

For **time** coordinate, a single axis is defined for a global referential: the T axis (in yellow, #F6DE2D), for which 1.0 corresponds to 1 second.

Other Conventions

- for clarity, we tend to use longer variable names, in CamelCase
- we tend to use mute variables to clarify meanings and intents, as in `_Acc=[]` (beware, despite being muted, any variable in scope that bears the same name will be matched), `Acc` designating accumulators
- as there is much list-based recursion, a variable named H means **Head** and T means **Tail** (as in `[Head|Tail]`)
- the string format specifier `~s` shall never be used; its Unicode-aware counterpart `~ts` must be used instead; similarly, for string operations, `list_to_binary/1` and `binary_to_list/1` must not be used either; prefer anyway the primitives in `text_utils`

Myriad-related Troubleshooting

Header/Module Dependencies

Only a very basic dependency between header files (`*.hrl`) and implementation files (`*.erl`) is managed.

As expected, if `X.hrl` changed, `X.beam` will be recompiled whether or not `X.erl` changed. However, any `Y.erl` that would include `X.hrl` would not be automatically recompiled.

Typically, when in doubt after having modified a record in a header file, just run `make rebuild` from the root of that layer (build is fast anyway, as quite parallel).

Third-Party Dependencies

Let's suppose we have an application named `Foo` that relies on Myriad. `Foo` may define additional dependencies, which may be:

- either mandatory or optional
- needed starting from build-time (ex: if relying on their headers and/or modules - including parse-transforms), or only at runtime

For a given **optional** dependency (ex: regarding JSON), a USE make variable is to be defined in the layer that introduced this dependency (ex: `USE_JSON`, introduced by Myriad, therefore to be listed in its `GNUmakevars.inc`). This variable allows to have our native build system register the associated additional include and ebin directories.

The first step to enable such a dependency (ex: the JSON support) is to set its USE variable to `true` (ex: `USE_JSON = true`), as it is expected to be disabled by default. Depending on the service at hand, a specific backend may have also to be selected (ex: either `USE JSX = true` or `USE JIFFY = true` to select a suitable JSON parser).

Finally, some supports may depend on others (ex: enabling `USE_REST` will enable `USE_JSON` in turn).

Runtime-only Third-Party Dependencies

The dependencies discussed here have to be found only when *executing* one's application; they can be installed:

- either manually, in which case the location of their `ebin` directory (typically an absolute path then) shall be specified in the code path (see, in `GNUmakevars.inc`, the `JSX_SOFTWARE_BASE` make variable for an example)
- or thanks to rebar, in which case they shall obey the same rules as the [Build-time Third-Party Dependencies](#) discussed below

Build-time Third-Party Dependencies

Myriad does not have such dependencies, but layers above in the software stack (like a layer that would be named `Foo`) may.

To have such dependencies (ex: let's suppose that the `jsx` JSON parser defined header files that one wants to include) *installed* as well when building one's project (ex: `Foo`), one may rely on `rebar`, and list them in the project's `foo/conf/rebar.config.template` file (ex: `{deps, [bar, jsx]}`.) from which the actual `rebar.config` is to be generated (use the `make set-rebar-conf` target for that).

The actual compilation will be done by our native build system in all cases, either directly (when running `make all`) or when using `rebar compile` (`rebar` hooks will then ensure that in practice the application is compiled with our native rules anyway). Therefore appropriate make variables (ex: `JSX_REBAR_BASE`, in `myriad/GNUMakevars.inc`) shall be defined so that the corresponding BEAM files installed through `rebar` can be found in this native context as well (through the `BEAM_DIRS` make variable).

Finally, such dependencies may or may not be listed in the `deps` entry of the `conf/foo.app.src` file¹⁹, depending on whether they are optional or not.

Myriad-level Third-Party Dependencies

Myriad as such has no mandatory dependency (except Erlang itself of course), but *optional* ones may be enabled, for:

- a basic [JSON](#) support (see our [json_utils](#) module), thanks to a suitable actual JSON parser, either [jsx](#) or [jiffy](#); note that the detection and use of these parsers are done transparently at runtime, hence none of them is a declared dependency of Myriad, which will adapt to any choice made by the overall application that includes both Myriad and one of such parsers (provided, as mentioned above, that the proper `USE_*` make variables are set)
- a first-level support of the [HDF5](#) file format (see our [hdf5_support](#) module), based on - and thus requiring - the [enhanced fork](#) that we made of [erlhdf5](#)
- [Python](#) (see our [python_utils](#) module), thanks to [erlport](#)
- [SQLite](#) (see our [sql_support](#) module), thanks to the SQLite 3 Erlang binding that we retained, [erlang-sqlite3](#)

As an example, let's suppose that we need a JSON support and that we want to rely on the `jsx` parser (our default choice) for that.

If applying our conventions, supposing that Erlang and Rebar3 are already installed (otherwise refer to the getting Erlang and [getting Rebar3](#) sections), `jsx` may be installed with:

¹⁹After having edited this file, run `make rebar3-create-app-file` afterwards in order to have the three other versions of it properly generated (namely `./_build/lib/foo/ebin/foo.app`, `ebin/foo.app` and `src/foo.app.src`).


```
$ mkdir -p ~/Software/jsx
$ cd ~/Software/jsx
$ git clone https://github.com/talentdeficit/jsx.git
$ ln -s jsx jsx-current-install
$ cd jsx/
$ rebar3 compile && rebar3 eunit
```

About the table module

This is a pseudo module, which is not meant to exist as such (no `table.erl`, no `table.beam`).

The Myriad parse transform replaces references to the `table` module by (generally) references to the `map_hashtable` module. See [table transformations](#) for more information.

Enabling the Interconnection of Erlang nodes

This is not a Myriad gotcha per se, but rather an Erlang one. Way too often, for obscure reasons Erlang nodes fail to connect to each other (especially with long names), and little to no information is available to diagnose the issue.

Safety Measures

In order to maximise the chances that nodes are able to `net_adm:ping/1` successfully each other:

- at least for testing, run VMs spawned with preferably the same **version** of Erlang
- ensure that they rely on the same **EPMD** (TCP) port (default Erlang one is 4369, while Myriad default one is 4506); check for example that all launched nodes of interest can be seen with: `epmd -port 4506 -names`
- check that they use the same **cookie**, either from the start (use the `-setcookie MY_COOKIE` command-line option) or after having changed it after the VM was launched
- ensure that no **firewall** gets in the way; one may take inspiration for example from our [iptables.rules-FullDisabling.sh](#) script
- finally check that the local **DNS resolution** complies with the surprisingly picky constraints demanded by the Erlang VM

For this last point, `/etc/hosts` is often the scene of the disaster. If your hostname is `hurricane` and your domain is `foobar.org`, then a line like the following one is known to work (whereas many variations of it may be deemed "incorrect"):

```
127.0.0.1 hurricane.foobar.org hurricane localhost.localdomain localhost
```

provided of course that, still in that file, you have not also a declaration such as:

```
192.168.0.5 hurricane.foobar.org hurricane
```

(setting one's IP shall better be done in one's profile in `/etc/netctl`, right?)

Testing & Troubleshooting

In order to **quick-check** whether long-name connectivity is available and to rule out the most obvious culprits, open two terminals.

In the first:

```
# Check (with root permissions) that the firewall rules are safe; for example:
$ iptables -nL
Chain INPUT (policy ACCEPT)
target      prot opt source                destination

Chain FORWARD (policy ACCEPT)
target      prot opt source                destination

Chain OUTPUT (policy ACCEPT)
target      prot opt source                destination

# Just to be on the safer side for this test:
$ killall beam.smp epmd

# Then launch the target first node:
$ ERL_EPMD_PORT=4032 erl -name n1 -setcookie aa
Erlang/OTP 23 [erts-11.1.4] [source] [64-bit] [smp:4:4] [ds:4:4:10] [async-threads:1]

Eshell V11.1.4 (abort with ^G)
(n1@hurricane.foobar.org)1>
```

In the second terminal, try to find the previous node:

```
$ ERL_EPMD_PORT=4032 erl -name n2 -setcookie aa
Erlang/OTP 23 [erts-11.1.4] [source] [64-bit] [smp:4:4] [ds:4:4:10] [async-threads:1]

Eshell V11.1.4 (abort with ^G)
(n2@hurricane.foobar.org)1> net_adm:ping('n1@hurricane.foobar.org').
pong
```

If you see **pong** here, run to the nearest altar and make a sacrifice to any Distribution God you may believe in (Norse ones being presumably the most effective here), and apply the hints listed in the [Enabling the Interconnection of Erlang nodes](#) section.

Support for Myriad

So you respected the [prerequisites](#) and [build](#) sections, and something went wrong? Generally we made sure that any detected error is blocking and loudly reported, with as much context as possible.

The simpler solution is then to [create a relevant issue](#).

For all other needs, please drop an email to the address listed on top of document. We do our best to answer on a timely basis.

Finally, provided that they meet licensing terms, scope and quality standards, contributions of all sorts are very welcome, be them porting efforts, increased test coverage, functional enrichments, documentation improvements, code enhancements, etc. See the next section for additional guidelines.

Please React!

If you have information more detailed or more recent than those presented in this document, if you noticed errors, neglects or points insufficiently discussed, drop us a line! (for that, follow the [Support](#) guidelines).

Contributions & Ending Word

Each time that you need a basic service that:

- seems neither provided by the Erlang [built-in modules](#) nor by this Myriad layer
- is generic-enough, simple and requires no special prerequisite

please either enrich our `*_utils.erl` [helpers](#), or add new general services!
In such a case, we would prefer that, in contributed code, the Myriad [Text Conventions](#) and [Coding Practices](#) are respected.

Thanks in advance, and have fun with Ceylan-Myriad!

The logo for MYRIAD, featuring the word "MYRIAD" in a stylized, blue, blocky font. Each letter is composed of multiple horizontal bars, giving it a striped or digital appearance.